Problem Sheet 1

Quantitative Methods

Preparation: Work through chapters 2 and 3 of Grolemund's *Hands-On Programming with R* (available here) and chapters 3–5 of Wickham and Grolemund's *R for Data Science* (available here). Do not just read — type in all the commands yourself, replicate the examples whenever possible, and try the exercises.

A. You roll a three dice and record the sum of the number of eyes on each die.

1. Use the `expand.grid()` function in R to create a data frame `S` that shows all possible outcomes of the experiment (the sample space).
2. Note that `expand.grid()` has automatically created three variables called `Var1`, `Var2`, and `Var3`. Create a new variable called `value` that records the sum of `Var1`, `Var2`, and `Var3`.
3. Assuming the dice are fair, calculate the probability that the sum of eyes is equal to 12.
4. Now assume each die is biased, with a probability of rolling a 6 equal to $\frac{3}{8}$ and the probability of all other outcomes equal to $\frac{1}{8}$. Create a vector recording these probabilities and assign them to the relevant entries in `S`.
5. Calculate the updated probability that the sum of eyes is equal to 12.

B.[1] Download the following CSV files and save them in your working directory: `Kenya.csv`, `Sweden.csv`, and `World.csv`. These files contain the following variables:

- ⋆ `country`: abbreviated country name.
- ⋆ `period`: period during which data are collected.
- ⋆ `age`: age group.
- ⋆ `births`: number of children (in thousands) born to women in each age group.
- ⋆ `deaths`: number of deaths (in thousands).
- ⋆ `py.men`: person-years for men (in thousands).
- ⋆ `py.women`: person-years for women (in thousands).

[1] This exercise draws on Exercise 1.5.2 in Kosuke Imai's *Quantitative Social Science: An Introduction*, Chapter 1. The relevant materials are also available here.

The data are collected for a period of 5 years, where *person-year* is a measure of the time contribution of each person during the period. For example, a person who lives through the entire 5-year period contributes 5 person-years, whereas a person who dies after 2 years contributes only 2 person-years.

1. Read each data set into R using either `read.csv()` or `tidyverse::read_csv()`.

2. Use the functions `summary()`, `head()`, and `tail()` to inspect each data set. You can also look directly at the data via `print()`, including the argument `n = 30` to see all 30 rows, or by double-clicking on the data frame in the `Environment` tab in RStudio.

3. The *age-specific fertility rate* (ASFR) within an age range $[x,\ x+n)$, where x is the starting age and n is the width of the age range (in years), is defined as

$$\text{ASFR}_{[x,\,x+n)} = \frac{\text{number of births to women of age } [x,\ x+n)}{\text{number of person-years lived by women of age } [x,\ x+n)}.$$

   Create a function called `asfr()` that computes the ASFR for each age group for women within the reproductive age range [15, 50). Calculate the ASFR for Kenya, Sweden, and the whole world separately for each of the two time periods in the data.

4. Using the ASFR, the *total fertility rate* (TFR) is defined as the average number of children given birth to by women who live through their entire reproductive age:

$$\text{TFR} = 5 \times (\text{ASFR}_{[15,\,19)} + \text{ASFR}_{[20,\,24)} + \cdots + \text{ASFR}_{[45,\,49)}).$$

   We multiply the sum by 5 because each woman spends five years in each age range, during which time her annual fertility rate is the ASFR. Create a function `tfr()` to compute the TFR for Kenya, Sweden, and the whole world separately for each of the two time periods in the data.

5. Now calculate the *age-specific death rate* (ASDR), which is defined as

$$\text{ASDR}_{[x,\,x+n)} = \frac{\text{number of deaths of people of age } [x,\ x+n)}{\text{number of person-years of people of age } [x,\ x+n)}.$$

   Create a function called `asdr()` to calculate the ASDR seperately by geography, age group, and time period.

6. Use `ggplot()` to visualise the the ASFR and the ASDR by geography, age group, and time period. Briefly summarise your principal findings.

Deadline: Submit a tidy and annotated R script via email by 2PM on Wednesday 21 October.