

Quantitative Methods

Human Sciences, 2020–21

Elias Nosrati

Lecture 3: 29 October 2020

Today

Today

- ▶ Recap on conditional probability: prosecutor's fallacy and the Monty Hall problem.

Today

- ▶ Recap on conditional probability: prosecutor's fallacy and the Monty Hall problem.
- ▶ Random variables and their distributions.

Today

- ▶ Recap on conditional probability: prosecutor's fallacy and the Monty Hall problem.
- ▶ Random variables and their distributions.
- ▶ Problem sheet 2 (tutorial).

Prosecutor's fallacy

Prosecutor's fallacy

- ▶ In 1998, Sally Clark was tried for murder after two of her sons died shortly after birth.

Prosecutor's fallacy

- ▶ In 1998, Sally Clark was tried for murder after two of her sons died shortly after birth.
- ▶ An expert witness for the prosecution testified that the probability of a newborn dying of sudden infant death syndrome (SIDS) was $1/8500$. Hence the probability of two consecutive deaths from SIDS was $(1/8500)^2$: ca. 1 in 73 million.

Prosecutor's fallacy

- ▶ In 1998, Sally Clark was tried for murder after two of her sons died shortly after birth.
- ▶ An expert witness for the prosecution testified that the probability of a newborn dying of sudden infant death syndrome (SIDS) was $1/8500$. Hence the probability of two consecutive deaths from SIDS was $(1/8500)^2$: ca. 1 in 73 million.
- ▶ The concluding argument: the probability of Clark's innocence was 1 in 73 million.

Prosecutor's fallacy

- ▶ In 1998, Sally Clark was tried for murder after two of her sons died shortly after birth.
- ▶ An expert witness for the prosecution testified that the probability of a newborn dying of sudden infant death syndrome (SIDS) was $1/8500$. Hence the probability of two consecutive deaths from SIDS was $(1/8500)^2$: ca. 1 in 73 million.
- ▶ The concluding argument: the probability of Clark's innocence was 1 in 73 million.
- ▶ What (if anything) is wrong with this line of reasoning?

Prosecutor's fallacy (cont.)

- ▶ Two problems:

Prosecutor's fallacy (cont.)

- ▶ Two problems:
 1. Independence assumption: how do we know that the two deaths are independent?

Prosecutor's fallacy (cont.)

- ▶ Two problems:
 1. Independence assumption: how do we know that the two deaths are independent?
 2. $\mathbb{P}(\text{innocence} \mid \text{evidence}) \neq \mathbb{P}(\text{evidence} \mid \text{innocence})$.

Prosecutor's fallacy (cont.)

- ▶ Two problems:
 1. Independence assumption: how do we know that the two deaths are independent?
 2. $\mathbb{P}(\text{innocence} \mid \text{evidence}) \neq \mathbb{P}(\text{evidence} \mid \text{innocence})$.
- ▶ By Bayes' Rule:

$$\mathbb{P}(\text{innocence} \mid \text{evidence}) = \frac{\mathbb{P}(\text{evidence} \mid \text{innocence})\mathbb{P}(\text{innocence})}{\mathbb{P}(\text{evidence})}.$$

Prosecutor's fallacy (cont.)

- ▶ Two problems:
 1. Independence assumption: how do we know that the two deaths are independent?
 2. $\mathbb{P}(\text{innocence} \mid \text{evidence}) \neq \mathbb{P}(\text{evidence} \mid \text{innocence})$.

- ▶ By Bayes' Rule:

$$\mathbb{P}(\text{innocence} \mid \text{evidence}) = \frac{\mathbb{P}(\text{evidence} \mid \text{innocence})\mathbb{P}(\text{innocence})}{\mathbb{P}(\text{evidence})}.$$

- ▶ Note that the prior probability of innocence is extremely high: double deaths due to SIDS might be rare but so are double infanticides!

Prosecutor's fallacy (cont.)

- ▶ Two problems:
 1. Independence assumption: how do we know that the two deaths are independent?
 2. $\mathbb{P}(\text{innocence} \mid \text{evidence}) \neq \mathbb{P}(\text{evidence} \mid \text{innocence})$.

- ▶ By Bayes' Rule:

$$\mathbb{P}(\text{innocence} \mid \text{evidence}) = \frac{\mathbb{P}(\text{evidence} \mid \text{innocence})\mathbb{P}(\text{innocence})}{\mathbb{P}(\text{evidence})}.$$

- ▶ Note that the prior probability of innocence is extremely high: double deaths due to SIDS might be rare but so are double infanticides!
- ▶ Re-write $\mathbb{P}(\text{evidence})$ using Law of Total Probability:
$$\mathbb{P}(\text{evidence} \mid \text{innocence})\mathbb{P}(\text{innocence}) + \mathbb{P}(\text{evidence} \mid \text{guilt})\mathbb{P}(\text{guilt}).$$

Prosecutor's fallacy (cont.)

- ▶ Two problems:
 1. Independence assumption: how do we know that the two deaths are independent?
 2. $\mathbb{P}(\text{innocence} \mid \text{evidence}) \neq \mathbb{P}(\text{evidence} \mid \text{innocence})$.

- ▶ By Bayes' Rule:

$$\mathbb{P}(\text{innocence} \mid \text{evidence}) = \frac{\mathbb{P}(\text{evidence} \mid \text{innocence})\mathbb{P}(\text{innocence})}{\mathbb{P}(\text{evidence})}.$$

- ▶ Note that the prior probability of innocence is extremely high: double deaths due to SIDS might be rare but so are double infanticides!
- ▶ Re-write $\mathbb{P}(\text{evidence})$ using Law of Total Probability:
$$\mathbb{P}(\text{evidence} \mid \text{innocence})\mathbb{P}(\text{innocence}) + \mathbb{P}(\text{evidence} \mid \text{guilt})\mathbb{P}(\text{guilt}).$$
- ▶ Note that $\mathbb{P}(\text{guilt}) \approx 0$

Prosecutor's fallacy (cont.)

- ▶ Two problems:
 1. Independence assumption: how do we know that the two deaths are independent?
 2. $\mathbb{P}(\text{innocence} \mid \text{evidence}) \neq \mathbb{P}(\text{evidence} \mid \text{innocence})$.

- ▶ By Bayes' Rule:

$$\mathbb{P}(\text{innocence} \mid \text{evidence}) = \frac{\mathbb{P}(\text{evidence} \mid \text{innocence})\mathbb{P}(\text{innocence})}{\mathbb{P}(\text{evidence})}.$$

- ▶ Note that the prior probability of innocence is extremely high: double deaths due to SIDS might be rare but so are double infanticides!

- ▶ Re-write $\mathbb{P}(\text{evidence})$ using Law of Total Probability:

$$\mathbb{P}(\text{evidence} \mid \text{innocence})\mathbb{P}(\text{innocence}) + \mathbb{P}(\text{evidence} \mid \text{guilt})\mathbb{P}(\text{guilt}).$$

- ▶ Note that $\mathbb{P}(\text{guilt}) \approx 0$ and hence

$$\begin{aligned}\mathbb{P}(\text{innocence} \mid \text{evidence}) &\approx \frac{\mathbb{P}(\text{evidence} \mid \text{innocence})\mathbb{P}(\text{innocence})}{\mathbb{P}(\text{evidence} \mid \text{innocence})\mathbb{P}(\text{innocence})} \\ &= 1.\end{aligned}$$

The Monty Hall problem

The Monty Hall problem

- ▶ Choose one of three doors: two with goats, one with car.

The Monty Hall problem

- ▶ Choose one of three doors: two with goats, one with car.
- ▶ After your choice, Monty Hall opens one door with a goat and asks if you want to switch to the remaining door. Should you switch?

The Monty Hall problem

- ▶ Choose one of three doors: two with goats, one with car.
- ▶ After your choice, Monty Hall opens one door with a goat and asks if you want to switch to the remaining door. Should you switch?
- ▶ Solve analytically or by simulation.

The Monty Hall problem: solving by simulation

```
sims <- 10 ^ 5 # Simulations  
doors <- 1:3 # Label doors  
win_no_switch <- 0 # Win count without switch  
win_switch <- 0 # Win count with switch
```

The Monty Hall problem: solving by simulation

```
sims <- 10 ^ 5 # Simulations
doors <- 1:3 # Label doors
win_no_switch <- 0 # Win count without switch
win_switch <- 0 # Win count with switch
```

```
for (i in 1:sims) {
  # Pick a door for the car
  car_door <- sample(doors, 1)
  # Pick a door for the contestant
  chosen_door <- sample(doors, 1)
  # If they match, add to "no switch" count
  if (car_door == chosen_door)
    win_no_switch <- win_no_switch + 1
  # Otherwise, add to switch count
  else
    win_switch <- win_switch + 1
}
```


The Monty Hall problem: solving by simulation (cont.)

```
cat("P(car | no switch) =", win_no_switch / sims)  
## P(car | no switch) = 0.33102
```

The Monty Hall problem: solving by simulation (cont.)

```
cat("P(car | no switch) =", win_no_switch / sims)  
## P(car | no switch) = 0.33102
```

```
cat("P(car | switch) =", win_switch / sims)  
## P(car | switch) = 0.66898
```

How is this possible?

The Monty Hall problem: analytic solution

The Monty Hall problem: analytic solution

- ▶ Conditioning as wishful thinking: condition on what you wish you knew.

The Monty Hall problem: analytic solution

- ▶ Conditioning as wishful thinking: condition on what you wish you knew.
- ▶ Let C_i be the event that the car is behind door i , for $1 \leq i \leq 3$, and let W be the event that you win the car.

The Monty Hall problem: analytic solution

- ▶ Conditioning as wishful thinking: condition on what you wish you knew.
- ▶ Let C_i be the event that the car is behind door i , for $1 \leq i \leq 3$, and let W be the event that you win the car.
- ▶ Since the C_i partition the sample space, we can use the Law of Total Probability:

$$\mathbb{P}(W) = \mathbb{P}(W \mid C_1) \times \frac{1}{3} + \mathbb{P}(W \mid C_2) \times \frac{1}{3} + \mathbb{P}(W \mid C_3) \times \frac{1}{3}.$$

The Monty Hall problem: analytic solution

- ▶ Conditioning as wishful thinking: condition on what you wish you knew.
- ▶ Let C_i be the event that the car is behind door i , for $1 \leq i \leq 3$, and let W be the event that you win the car.
- ▶ Since the C_i partition the sample space, we can use the Law of Total Probability:

$$\mathbb{P}(W) = \mathbb{P}(W \mid C_1) \times \frac{1}{3} + \mathbb{P}(W \mid C_2) \times \frac{1}{3} + \mathbb{P}(W \mid C_3) \times \frac{1}{3}.$$

- ▶ Suppose, without loss of generality, that you have chosen door 1 but decide to switch.

The Monty Hall problem: analytic solution

- ▶ Conditioning as wishful thinking: condition on what you wish you knew.
- ▶ Let C_i be the event that the car is behind door i , for $1 \leq i \leq 3$, and let W be the event that you win the car.
- ▶ Since the C_i partition the sample space, we can use the Law of Total Probability:

$$\mathbb{P}(W) = \mathbb{P}(W \mid C_1) \times \frac{1}{3} + \mathbb{P}(W \mid C_2) \times \frac{1}{3} + \mathbb{P}(W \mid C_3) \times \frac{1}{3}.$$

- ▶ Suppose, without loss of generality, that you have chosen door 1 but decide to switch.
- ▶ If C_1 , the switching strategy will fail:

$$\mathbb{P}(W \mid C_1) = 0.$$

The Monty Hall problem: analytic solution

- ▶ Conditioning as wishful thinking: condition on what you wish you knew.
- ▶ Let C_i be the event that the car is behind door i , for $1 \leq i \leq 3$, and let W be the event that you win the car.
- ▶ Since the C_i partition the sample space, we can use the Law of Total Probability:

$$\mathbb{P}(W) = \mathbb{P}(W \mid C_1) \times \frac{1}{3} + \mathbb{P}(W \mid C_2) \times \frac{1}{3} + \mathbb{P}(W \mid C_3) \times \frac{1}{3}.$$

- ▶ Suppose, without loss of generality, that you have chosen door 1 but decide to switch.
- ▶ If C_1 , the switching strategy will fail:

$$\mathbb{P}(W \mid C_1) = 0.$$

- ▶ If the car is behind doors 2 or 3, switching will succeed:

$$\mathbb{P}(W \mid C_2) = \mathbb{P}(W \mid C_3) = 1.$$

The Monty Hall problem: analytic solution

- ▶ Conditioning as wishful thinking: condition on what you wish you knew.
- ▶ Let C_i be the event that the car is behind door i , for $1 \leq i \leq 3$, and let W be the event that you win the car.
- ▶ Since the C_i partition the sample space, we can use the Law of Total Probability:

$$\mathbb{P}(W) = \mathbb{P}(W \mid C_1) \times \frac{1}{3} + \mathbb{P}(W \mid C_2) \times \frac{1}{3} + \mathbb{P}(W \mid C_3) \times \frac{1}{3}.$$

- ▶ Suppose, without loss of generality, that you have chosen door 1 but decide to switch.
- ▶ If C_1 , the switching strategy will fail:

$$\mathbb{P}(W \mid C_1) = 0.$$

- ▶ If the car is behind doors 2 or 3, switching will succeed:

$$\mathbb{P}(W \mid C_2) = \mathbb{P}(W \mid C_3) = 1.$$

- ▶ Hence, $\mathbb{P}(W) = 0 \times \frac{1}{3} + 1 \times \frac{1}{3} + 1 \times \frac{1}{3} = \frac{2}{3}$.

Random variables

Random variables

Definition

Given an experiment with sample space S , a *random variable* is a function from the sample space S to the set of real numbers \mathbb{R} . Thus a random variable X assigns a numerical value $X(s)$ to each possible outcome s of the experiment:

$$X : S \rightarrow \mathbb{R}.$$

Random variables

Definition

Given an experiment with sample space S , a *random variable* is a function from the sample space S to the set of real numbers \mathbb{R} . Thus a random variable X assigns a numerical value $X(s)$ to each possible outcome s of the experiment:

$$X : S \rightarrow \mathbb{R}.$$

- ▶ A concept that simplifies notation.

Random variables

Definition

Given an experiment with sample space S , a *random variable* is a function from the sample space S to the set of real numbers \mathbb{R} . Thus a random variable X assigns a numerical value $X(s)$ to each possible outcome s of the experiment:

$$X : S \rightarrow \mathbb{R}.$$

- ▶ A concept that simplifies notation.
- ▶ Expands our ability to quantify uncertainty and summarise the results of experiments.

Random variables

Definition

Given an experiment with sample space S , a *random variable* is a function from the sample space S to the set of real numbers \mathbb{R} . Thus a random variable X assigns a numerical value $X(s)$ to each possible outcome s of the experiment:

$$X : S \rightarrow \mathbb{R}.$$

- ▶ A concept that simplifies notation.
- ▶ Expands our ability to quantify uncertainty and summarise the results of experiments.
- ▶ The function itself is deterministic.

Random variables


Definition

Given an experiment with sample space S , a *random variable* is a function from the sample space S to the set of real numbers \mathbb{R} . Thus a random variable X assigns a numerical value $X(s)$ to each possible outcome s of the experiment:

$$X : S \rightarrow \mathbb{R}.$$

- ▶ A concept that simplifies notation.
- ▶ Expands our ability to quantify uncertainty and summarise the results of experiments.
- ▶ The function itself is deterministic.
- ▶ The randomness derives from the experiment whose outcomes have probabilities described by the probability function $\mathbb{P}(\cdot)$.

Random variables (cont.)



Random variables: coin flip example

Random variables: coin flip example

- ▶ We flip a fair coin twice.

Random variables: coin flip example

- ▶ We flip a fair coin twice.
- ▶ $S = \{HH, HT, TH, TT\}$.

Random variables: coin flip example

- ▶ We flip a fair coin twice.
- ▶ $S = \{HH, HT, TH, TT\}$.
- ▶ Let X be the number of Heads. This is a random variable with possible values 0, 1, and 2:

$$X(HH) = 2, \quad X(HT) = X(TH) = 1, \quad X(TT) = 0.$$

Random variables: coin flip example

- ▶ We flip a fair coin twice.
- ▶ $S = \{HH, HT, TH, TT\}$.
- ▶ Let X be the number of Heads. This is a random variable with possible values 0, 1, and 2:

$$X(HH) = 2, \quad X(HT) = X(TH) = 1, \quad X(TT) = 0.$$

- ▶ Let Y be the number of tails. Then $Y = 2 - X$.

Random variables: coin flip example

- ▶ We flip a fair coin twice.
- ▶ $S = \{HH, HT, TH, TT\}$.
- ▶ Let X be the number of Heads. This is a random variable with possible values 0, 1, and 2:

$$X(HH) = 2, \quad X(HT) = X(TH) = 1, \quad X(TT) = 0.$$

- ▶ Let Y be the number of tails. Then $Y = 2 - X$.
- ▶ Let I be the *indicator random variable* that is 1 if the first flip is Heads and 0 otherwise.

Random variables: coin flip example

- ▶ We flip a fair coin twice.
- ▶ $S = \{HH, HT, TH, TT\}$.
- ▶ Let X be the number of Heads. This is a random variable with possible values 0, 1, and 2:

$$X(HH) = 2, \quad X(HT) = X(TH) = 1, \quad X(TT) = 0.$$

- ▶ Let Y be the number of tails. Then $Y = 2 - X$.
- ▶ Let I be the *indicator random variable* that is 1 if the first flip is Heads and 0 otherwise.
- ▶ Note that two random variables can be defined on the same sample space: the outcomes are the same, but the numerical values assigned to the outcomes are different.

Discrete random variables

Discrete random variables

Definition

A random variable X is said to be *discrete* if it takes on a finite (or at most countably infinite) number of values. In other words, its outcomes are discrete.

Discrete random variables

Definition

A random variable X is said to be *discrete* if it takes on a finite (or at most countably infinite) number of values. In other words, its outcomes are discrete.

Example

The number of Heads resulting from a single coin flip is a *binary* random variable: it takes on only two values (0 or 1). The sum of the number of eyes on two dice or the number of votes won by a presidential candidate are also discrete random variables.

Continuous random variables

Continuous random variables

Definition

A random variable X is said to be *continuous* if it takes on an (uncountably) infinite number of values. In other words, it may take on a continuum of possible values.

Continuous random variables

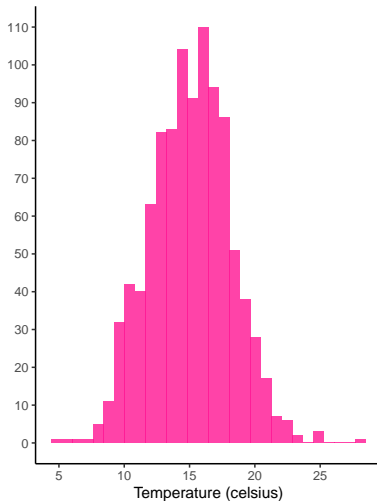

Definition

A random variable X is said to be *continuous* if it takes on an (uncountably) infinite number of values. In other words, it may take on a continuum of possible values.

Example

The total amount of rainfall over the next year is a continuous random variable. So is the temperature on a randomly chosen day, or the height of a randomly selected person.

The distribution of random variables



Probability distributions

Probability distributions

Definition

The *probability distribution* of a random variable X is the assignment of a probability measure to every conceivable realisation of X .

Probability distributions

Definition

The *probability distribution* of a random variable X is the assignment of a probability measure to every conceivable realisation of X .

- ▶ The distribution of a random variable X specifies the probabilities of all events associated with X .

Probability distributions

Definition

The *probability distribution* of a random variable X is the assignment of a probability measure to every conceivable realisation of X .

- ▶ The distribution of a random variable X specifies the probabilities of all events associated with X .
- ▶ Describes which values of X are more likely than other values.

Probability distributions

Definition

The *probability distribution* of a random variable X is the assignment of a probability measure to every conceivable realisation of X .

- ▶ The distribution of a random variable X specifies the probabilities of all events associated with X .
- ▶ Describes which values of X are more likely than other values.
- ▶ Commonly encoded by a *probability density function*.

Probability distributions

Definition

The *probability distribution* of a random variable X is the assignment of a probability measure to every conceivable realisation of X .

- ▶ The distribution of a random variable X specifies the probabilities of all events associated with X .
- ▶ Describes which values of X are more likely than other values.
- ▶ Commonly encoded by a *probability density function*.
- ▶ Notation: $f(x) = f(X = x)$.

Probability distributions

Definition

The *probability distribution* of a random variable X is the assignment of a probability measure to every conceivable realisation of X .

- ▶ The distribution of a random variable X specifies the probabilities of all events associated with X .
- ▶ Describes which values of X are more likely than other values.
- ▶ Commonly encoded by a *probability density function*.
- ▶ Notation: $f(x) = f(X = x)$.
- ▶ Two requirements:

Probability distributions

Definition

The *probability distribution* of a random variable X is the assignment of a probability measure to every conceivable realisation of X .

- ▶ The distribution of a random variable X specifies the probabilities of all events associated with X .
- ▶ Describes which values of X are more likely than other values.
- ▶ Commonly encoded by a *probability density function*.
- ▶ Notation: $f(x) = f(X = x)$.
- ▶ Two requirements:
 1. $f(x) \geq 0$ for any x .

Probability distributions

Definition

The *probability distribution* of a random variable X is the assignment of a probability measure to every conceivable realisation of X .

- ▶ The distribution of a random variable X specifies the probabilities of all events associated with X .
- ▶ Describes which values of X are more likely than other values.
- ▶ Commonly encoded by a *probability density function*.
- ▶ Notation: $f(x) = f(X = x)$.
- ▶ Two requirements:
 1. $f(x) \geq 0$ for any x .
 2. $\sum f(x) = 1$ (discrete) or $\int f(x) = 1$ (continuous).

What exactly is a density?

What exactly is a density?

- ▶ Let X be a random variable (e.g., age of random person).

What exactly is a density?



- ▶ Let X be a random variable (e.g., age of random person).
- ▶ To visualise the distribution of X , we 'discretise' the variable by creating bins or intervals along the variable of interest and count the number of observations that fall within each bin.

What exactly is a density?



- ▶ Let X be a random variable (e.g., age of random person).
- ▶ To visualise the distribution of X , we 'discretise' the variable by creating bins or intervals along the variable of interest and count the number of observations that fall within each bin.
- ▶ The *density* for each bin is then defined as

$$\text{density} = \frac{\text{proportion of observations in the bin}}{\text{width of the bin}}.$$

What exactly is a density? (cont.)





What exactly is a density? (cont.)





- We interpret the density scale as a probability (or percentage) per horizontal unit.

What exactly is a density? (cont.)



- ▶ We interpret the density scale as a probability (or percentage) per horizontal unit.
- ▶ For a discrete random variable, probability = density.

What exactly is a density? (cont.)



- ▶ We interpret the density scale as a probability (or percentage) per horizontal unit.
- ▶ For a discrete random variable, probability = density.
- ▶ For a continuous random variable, probability = area under curve.